

```
install.packages("ergm", dependencies = TRUE)
install.packages("network")
install.packages("parallel")
install.packages("gridExtra")
```

Installing package into ‘/usr/local/lib/R/site-library’
(as ‘lib’ is unspecified)

also installing the dependencies ‘ergm.count’, ‘Rglpk’

Warning message in install.packages("ergm", dependencies = TRUE):
“installation of package ‘Rglpk’ had non-zero exit status”
Installing package into ‘/usr/local/lib/R/site-library’
(as ‘lib’ is unspecified)

Installing package into ‘/usr/local/lib/R/site-library’
(as ‘lib’ is unspecified)

Warning message:
“package ‘parallel’ is a base package, and should not be updated”
Installing package into ‘/usr/local/lib/R/site-library’
(as ‘lib’ is unspecified)

```
library(ergm)
library(network)
library(parallel)
library(ggplot2)
library(gridExtra)
```

```
# Had to use multithreaded approach to run a reduced ERGM model.
num_cores <- detectCores()
```

```
# Method to plot coefficients from ERGM model.
plot_model_coefs <- function(model, title) {
  coefs <- summary(model)$coefficients
  df <- data.frame(term = rownames(coefs), estimate = coefs[, 1], se =
coefs[, 2])
  df$ci_lower <- df$estimate - 1.96 * df$se
  df$ci_upper <- df$estimate + 1.96 * df$se

  ggplot(df, aes(x = term, y = estimate)) +
    geom_point() +
    geom_errorbar(aes(ymin = ci_lower, ymax = ci_upper), width = 0.2)
+
  coord_flip() +
  labs(title = title, y = "Estimate", x = "") +
  theme_minimal()
}
```

```
## -- Load Caltech's graph --

# Load adjacency matrix (skip first row with node IDs)
caltech_adj <- as.matrix(read.csv("Caltech36_adj.csv", header = FALSE,
skip = 1))

# Load node attributes.
caltech_node_info <- read.csv("Caltech36_local_info.csv")

# Create undirected network.
caltech_g <- network(caltech_adj, directed = FALSE)

# Add node-level attributes.
set.vertex.attribute(caltech_g, "year", caltech_node_info$year)
set.vertex.attribute(caltech_g, "residence", caltech_node_info$dorm)
set.vertex.attribute(caltech_g, "major", caltech_node_info$major)
set.vertex.attribute(caltech_g, "high_school",
caltech_node_info$high_school)

# Run ERGM model with modelling assumptions about attributes.
caltech_model <- ergm(
  caltech_g ~ edges + nodematch("year") + nodematch("residence") +
  nodematch("major") + nodematch("high_school"),
  control = control.ergm(
    parallel = num_cores,
    parallel.type = "PSOCK",
    MCMLE.maxit = 100
  )
)

# Print summary about the model
summary(caltech_model)

# Plot coefficients.
caltech_p <- plot_model_coefs(caltech_model, "Caltech")

grid.arrange(caltech_p, nrow=1)

Starting maximum pseudolikelihood estimation (MPLE):

Obtaining the responsible dyads.

Evaluating the predictor and response matrix.

Maximizing the pseudolikelihood.

Finished MPLE.

Evaluating log-likelihood at the estimate.
```

```
Call:
ergm(formula = caltech_g ~ edges + nodematch("year") +
nodematch("residence") +
      nodematch("major") + nodematch("high_school"), control =
control.ergm(parallel = num_cores,
parallel.type = "PSOCK", MCMLE.maxit = 100))
```


Maximum Likelihood Results:

	Estimate	Std. Error	MCMC %	z value	Pr(> z)	
edges	-3.59328	0.01254	0	-286.47	<1e-04	***
nodematch.year	1.30579	0.01746	0	74.81	<1e-04	***
nodematch.residence	1.80352	0.01734	0	104.03	<1e-04	***
nodematch.major	0.37469	0.02797	0	13.39	<1e-04	***
nodematch.high_school	-2.07340	0.07712	0	-26.88	<1e-04	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Null Deviance: 409367 on 295296 degrees of freedom
Residual Deviance: 112140 on 295291 degrees of freedom

AIC: 112150 BIC: 112203 (Smaller is better. MC Std. Err. = 0)



```
## -- Load Dartmouth's graph --

# Load adjacency matrix (skip first row with node IDs)
dartmouth_adj <- as.matrix(read.csv("Dartmouth6_adj.csv", header =
FALSE, skip = 1))

# Load node attributes.
dartmouth_node_info <- read.csv("Dartmouth6_local_info.csv")

# Create undirected network.
dartmouth_g <- network(dartmouth_adj, directed = FALSE)

# Add node-level attributes.
set.vertex.attribute(dartmouth_g, "year", dartmouth_node_info$year)
set.vertex.attribute(dartmouth_g, "residence",
dartmouth_node_info$dorm)
set.vertex.attribute(dartmouth_g, "major", dartmouth_node_info$major)
set.vertex.attribute(dartmouth_g, "high_school",
dartmouth_node_info$high_school)

# Run ERGM model with modelling assumptions about attributes.
dartmouth_model <- ergm(
  dartmouth_g ~ edges + nodematch("year") + nodematch("residence") +
  nodematch("major") + nodematch("high_school"),
  control = control.ergm(
    parallel = num_cores,
    parallel.type = "PSOCK",
    MCMLE.maxit = 100
  )
)

# Print summary about the model
summary(dartmouth_model)

# Plot coefficients.
dartmouth_p <- plot_model_coefs(dartmouth_model, "Dartmouth")

grid.arrange(dartmouth_p, nrow=1)

Starting maximum pseudolikelihood estimation (MPLE):

Obtaining the responsible dyads.

Evaluating the predictor and response matrix.

Maximizing the pseudolikelihood.

Finished MPLE.

Evaluating log-likelihood at the estimate.
```

```
Call:
ergm(formula = dartmouth_g ~ edges + nodematch("year") +
      nodematch("residence") +
      nodematch("major") + nodematch("high_school"), control =
control.ergm(parallel = num_cores,
      parallel.type = "PSOCK", MCMLE.maxit = 100))
```

Maximum Likelihood Results:

	Estimate	Std. Error	MCMC %	z value	Pr(> z)
edges	-4.498292	0.002167	0	-2076.28	<1e-04

nodematch.year	0.102176	0.005597	0	18.26	<1e-04

nodematch.residence	-0.191384	0.005365	0	-35.67	<1e-04


nodematch.major	-0.149691	0.007036	0	-21.28	<1e-04

nodematch.high_school	-0.581930	0.007818	0	-74.44	<1e-04


Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Null Deviance: 40325423 on 29088644 degrees of freedom
Residual Deviance: 3298500 on 29088639 degrees of freedom


AIC: 3298510 BIC: 3298586 (Smaller is better. MC Std. Err. = 0)



```
# Plot side-by-side  
grid.arrange(caltech_p, dartmouth_p, nrow = 1)
```



```
# Plot side-by-side  
grid.arrange(caltech_p, dartmouth_p)
```



```
## -- Load Cornell's graph --
## Ran into RAM limitations for loading the adjacency matrix on my
MacBook.

# Load adjacency matrix (skip first row with node IDs)
cornell_adj <- as.matrix(read.csv("Cornell5_adj.csv", header = FALSE,
skip = 1))

# Load node attributes.
cornell_node_info <- read.csv("Cornell5_local_info.csv")

# Create undirected network.
cornell_g <- network(cornell_adj, directed = FALSE)

# Add node-level attributes.
set.vertex.attribute(cornell_g, "year", cornell_node_info$year)
set.vertex.attribute(cornell_g, "residence", cornell_node_info$dorm)
set.vertex.attribute(cornell_g, "major", cornell_node_info$major)
set.vertex.attribute(cornell_g, "high_school",
cornell_node_info$high_school)

# Run ERGM model with modelling assumptions about attributes.
cornell_model <- ergm(
  cornell_g ~ edges + nodematch("year") + nodematch("residence") +
```



```

    nodematch("major") + nodematch("high_school"),
    control = control.ergm(
      parallel = num_cores,
      parallel.type = "PSOCK",
      MCMLE.maxit = 100
    )
  )
)

# Print summary about the model
summary(cornell_model)

# Plot coefficients.
cornell_p <- plot_model_coefs(cornell_model, "Cornell")

grid.arrange(cornell_p, nrow=1)

```

Starting maximum pseudolikelihood estimation (MPLE):

Obtaining the responsible dyads.

Evaluating the predictor and response matrix.

Maximizing the pseudolikelihood.

Finished MPLE.

Evaluating log-likelihood at the estimate.

Call:

```

ergm(formula = cornell_g ~ edges + nodematch("year") +
  nodematch("residence") +
    nodematch("major") + nodematch("high_school"), control =
  control.ergm(parallel = num_cores,
    parallel.type = "PSOCK", MCMLE.maxit = 100))

```

Maximum Likelihood Results:

	Estimate	Std. Error	MCMC %	z value	Pr(> z)
edges	-5.342019	0.003227	0	-1655.514	<1e-04

nodematch.year	0.089152	0.007943	0	11.224	<1e-04

nodematch.residence	-0.228957	0.007902	0	-28.974	<1e-04

nodematch.major	-0.040502	0.013827	0	-2.929	0.0034
**					
nodematch.high_school	-0.423392	0.016145	0	-26.224	<1e-04


```

---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Null Deviance: 39940011 on 28810628 degrees of freedom
Residual Deviance: 1677430 on 28810623 degrees of freedom

AIC: 1677440 BIC: 1677516 (Smaller is better. MC Std. Err. = 0)

```




```


# Plot side-by-side
grid.arrange(caltech_p, dartmouth_p, cornell_p)

```

Caltech



Dartmouth



Cornell

